

Homework 2

Suppose I have a relation Grades(student_id, assignment_id, score). I have 200 students and 20 assignments. I would grade all submissions of one assignment based on the submission order, and then insert the records. As a result, based on my insertion nature, the student_id is not sorted, but the assignment_id is. I choose heap file as my file organization. My page is quite small – it can only store 40 records, or 200 bytes in one page. The SearchKeySize is 2 bytes and PointerSize is 2 bytes. My buffer size is also small, 4 pages.

1. (50 points) If my most frequent query is to find individual students, such as `select * from grades where student_id='3347';`
 - a. (5pts) What is the I/O cost (i.e., number of pages in terms of reading and writing) for this query if I don't build index for student_id? (note: student_id can appear as many as 20 times in this relation)

→ We have 200 students and 20 assignments.
Hence, the number of records can be calculated as:

$$\begin{aligned}\text{Number of records} &= \text{number of students} * \text{number of assignments} \\ &= 200 * 20 \\ &= 4000.\end{aligned}$$

Hence, we have 4000 records.

Now, the number of pages can be calculated as:

$$\begin{aligned}\text{Number of pages} &= \text{Number of records} / \text{Number of records a page can hold} \\ &= 4000 / 40 \\ &= 100.\end{aligned}$$

Therefore, the number of pages in terms of reading and writing for this query if we don't build and index for student_id is 100 pages, and it needs to read the entire relation.

- b. I want to improve the I/O cost. I am debating if I need to build index for student_id, or to sort based on student_id. So I need to do some estimation. Please help me by answering the following questions.

i. (15pts) What is the I/O cost of multi-way merge sort (aka, external sort) if I sort the relation after I enter all records? Explain the process.

➔ The first phase would be 'partitioning'. So, the I/O cost will be:

- i) For reading = 100
- ii) For writing = 100

Total: I/O = 200 (We will get 25 subfiles)

The second phase will include three rounds. In the first round, you will have to merge every 3 subfiles into one.

Then, the I/O cost would be: - For reading = 100

- For writing = 100

Total I/O = 200 . Now we will get 9 subfiles.

Now in the second round, you will again merge every 3 subfiles into one.

Then, the I/O cost would be: - For reading = 100

- For writing = 100

Total I/O = 200 . Now we will get 3 subfiles.

Now, for the final third round, we will again merge every 3 subfiles into one.

Then, the I/O cost would be: - For reading = 100

- For writing = 100

Total I/O = 200 .

■ **Here the process has been completed.**

And the final I/O cost of multi-way merge will be; 200+200+200+200 = 800.

ii. (15pts) Suppose I decide to build B+ tree index instead of sorting. What is the smallest number of pages do you estimate the B+ tree will take?

➔ Using the formula; $n * SearchKeySize + (n+1) * PointerSize \leq PageSize$

$$n * 2 + (n+1) * 2 \leq 200$$

$$2n + 2n + 2 \leq 200$$

$$4n \leq 198$$

Therefore, $n \leq 49.5$; $n = 49$ & $n+1 = 50$.

Hence, each page can store 49 keys and 50 pointers.

Now, as previously calculated, we have 4000 records. And in a dense index method, we have one pointer per record.

Then, on the leaf level in the B+ tree index, we will have: 4000 pointers
Which means, we will have, $4000 / 49 = \underline{82 \text{ pages on the leaf level.}}$

Now, the B+ tree index has three- four levels.

Hence, the smallest number of pages that the B+ tree will take is:

$$\underline{82 + 2 + 1 = 85 \text{ pages.}}$$

iii. (15pts) What is the worst I/O cost for answering this query with B+ tree index now?

→ We have 20 assignments to read, hence using the B+ tree index, we will have to read 20 data pages.

Now, the B+ tree index has three or four levels. So, the tree index will need/read three or four index pages.

Therefore, the worst I/O cost for answering this query with B+ tree index would be either;
 $20 + 3 = 23 \text{ pages or } 20 + 4 = 24 \text{ pages.}$

2. (40 points) If my most frequent query is to find all scores for an assignment, such as

select * from grades where assignment_id='01';

**a. (10pts) What is the I/O cost if I don't build index for assignment_id?
(note: assignment_id is sorted and each assignment_id can appear as many as 200 times in this relation)**

→ We have 100 pages.

Hence, for using the formula for binary search, we get;

$$\text{Log}_2 (100) = 6 \text{ or } 7 \text{ pages}$$

Now, a page can hold up to 40 records.

Therefore, to retrieve all the records; $200 / 40 = 5 \text{ pages.}$

Therefore, the total I/O cost if we don't build index for assignment id is;

$$\underline{6 + 5 = 11 \text{ pages or } 7 + 5 = 12 \text{ pages}}$$

- b. I am debating if building index for `assignment_id` would further improve the I/O cost. Please help me by answering the following questions.
- i. (15pts) Suppose I decide to build B+ tree index. What is the smallest number of pages do you estimate the B+ tree will take?

→ Now, in a sparse index method, we have one pointer per page in the B+ tree index. Since we calculated that we have 100 pages, On the leaf level of the B+ tree index, we will have 100 pointers.

Then, the numbers of pages on the leaf level will be; $100 / 49 = 3$ pages.

Here, the tree has two levels. So, the smallest number of pages the B+ tree will take is; $3 + 1 = 4$ pages.

- ii. (15pts) What is the best I/O cost for answering this query with B+ tree index now?

→ Here, the B+ tree index will read about 5 ± 1 data pages and it will read about two or three index pages.

Therefore,

The best I/O cost for answering this query with B+ tree index now will be either, 6, 7, 8, or 9.

3. (10 points) Suppose at the end of the semester, I need to curve the grades. I decide to increase all scores by 5 points. What is the I/O cost for this operation?

